

Hitesh Nandakumar

Dr. Riesen

PHIL 323

5/13/2026

Case Study 1: Autonomous Vehicles

I'll organize the data autonomous vehicles (AVs) can permissibly use into three tiers based on their source and ethical complexity. AVs can permissibly use all the data available to a traditional human driver, without any objection. This 1st degree data would include 360 degree video coverage, as well as information about the vehicle itself, such as weight and length. This raises no ethical objections, as if a human driver would reasonably have access to the data, then the AV by extension should also. The 2nd degree category comes from sensors, such as LIDAR and ultrasonic sensors. While this kind of information is not available to the human driver, its access serves to increase the safety of the passengers and others. It also does not pose any reasonable ethical objections that would outweigh its safety benefits. The last category of information is what I would consider 3rd degree information. This data comes from outside the car. Current examples are traffic data, accidents reported by users, and emergency vehicles nearby. In the future this could be expanded to access to traffic signal data and advanced metrics for efficient route planning and broader traffic mitigation. This kind of external data connection is being explored and raises some ethical questions, particularly around what kind of data the AVs should not have access to.

The easy rationale to justify the access of all pertinent tertiary data, such as passenger information in other cars and driving/accident history is that it can be used to mitigate harm in

an accident. I argue that this raises too many ethical objections to be allowed. By providing AVs with such 3rd degree information, it could potentially help reduce damage and harm, but on the other hand it enables these systems to make decisions beyond the reasonable capability of a human driver. For example, if an AV could access a pedestrian's medical records or criminal history, it might factor that into a crash decision, effectively passing judgement on someone's life based on data they never consented to share in the first place. Even though someone can be held responsible for the system's actions, it does not justify the system having the capacity to make these decisions. The data should be limited to information with the reach of a human driver and reasonable external information.

When a crash is inevitable, I argue that an AV can only make one decision: how to best ensure the safety of the passengers inside using reasonable actions. There are significant ethical challenges that arise when we attempt to distribute risks and harms at the hands of a machine. A civilian-use machine has no right to dictate who gets to live and who doesn't, even if the end goal is to minimize total impact. One might argue that a machine capable of minimizing total harm should do so; but this assumes a moral calculation that nobody has authorized the machine to perform. A human driver in the same situation would instinctively act to protect themselves, and we don't hold them morally responsible for that instinct. Expecting an AV to do otherwise, to coldly calculate whose life is worth more, assigns a kind of moral authority to a machine that no engineer, policymaker, or passenger can ever provide.